

How do we write parallel programs?

Wednesday, February 7, 2024

6:48 PM

- Task parallelism
 - Partition various tasks carried out solving the problem among the cores.
- Data parallelism
 - Partition the data used in solving the problem among the cores.
 - Each core carries out similar operations on it's part of the data.

Example of both:

Professor P: A problem I can relate to



15 questions
300 exams
3 TAs



Professor P's grading assistants




TA#1

TA#2

TA#3

Division of Work - data parallelism

Division of work – data parallelism



Division of work – data parallelism


```
sum = 0;
for (i = 0; i < n; i++) {
    x = Compute_next_value(. . .);
    sum += x;
}
```

Division of Work - task parallelism

TA#1



Division of work – task parallelism



Division of work – task parallelism

```

if (I'm the master core) {
    sum = my_x;
    for each core other than myself {
        receive value from core;
        sum += value;
    }
} else {
    send my_x to the master;
}
    
```

[Tasks](#)

1) [Receiving](#)
2) [Addition](#)

Coordination

- Cores usually need to coordinate their work.
- [Communication](#) – one or more cores send their current partial sums to another core.
- [Load balancing](#) – share the work evenly among the cores so that one is not heavily loaded.
- [Synchronization](#) – because each core works at its own pace, make sure cores do not get too far ahead of the rest.

Notes:

- Fundamentals of parallelism