

CUDA hello

The CUDA API

- API we'll be using to program heterogeneous CPU–GPU systems.
- CUDA is a software platform that can be used to write GPGPU programs for heterogeneous systems equipped with an Nvidia GPU
- CUDA was originally an acronym for "Compute Unified Device Architecture," which was meant to suggest that it provided a single interface for programming both CPU and GPU
 - o More recently, however, Nvidia has decided that CUDA is not an acronym; it's simply the name of an API for GPGPU programming
- There is a language-specific CUDA API for several languages; for example, there are CUDA APIs for C, C++, Fortran, Python, and Java.
 - o We'll be using CUDA C
 - but we need to be aware that sometimes we'll need to use some C++ constructs.
 - This is because the CUDA C compiler can compile both C and C++ programs, and it can do this because it is a modified C++ compiler.
 - So where the specifications for C and C++ differ the CUDA compiler sometimes uses C++. (preference for c++)
 - For example, since the C library function malloc returns a void* pointer
 - a C program doesn't need a cast in the instruction `float *x = malloc (n* sizeof (float));`
 - However, in C++ a cast is required `float *x = (float *) malloc (n* sizeof (float));`

As usual, we'll begin by implementing a version of the "hello, world" program.

- We'll write a CUDA C program in which each CUDA thread prints a greeting.
- Since the program is heterogeneous, we will, effectively, be writing two programs: a host or CPU program and a device or GPU program.
- Special compiler for CUDA
 - Note that even though our programs are written in CUDA C, CUDA programs cannot be compiled with an ordinary C compiler. So unlike MPI and Pthreads, CUDA is not just a library that can be linked in to an ordinary C program: CUDA requires a special compiler. For example, an ordinary C compiler (such as gcc) generates a machine language executable for a single CPU (e.g., an x86 processor), but the CUDA compiler must generate machine language for two different processors: the host processor and the device processor.
- Source Code

```
1 #include <stdio.h>
2 #include <cuda.h>  /* Header file for CUDA */
3
4 /* Device code: runs on GPU */
5 __global__ void Hello(void) {
6
7     printf("Hello from thread %d\n", threadIdx.x);
8 } /* Hello */
9
10 /* Host code: Runs on CPU */
11 int main(int argc, char* argv[]) {
12     int thread_count;    /* Number of threads to run on GPU */
13
14     thread_count = strtol(argv[1], NULL, 10);
15                     /* Get thread_count from command line */
16
17     Hello <<<1, thread_count>>>();
18                     /* Start thread_count threads on GPU, */
19
20     cudaDeviceSynchronize();    /* Wait for GPU to finish */
21
22     return 0;
23 } /* main */
24
```

Program 6.1: CUDA program that prints greetings from the threads.

- Output

```
Hello from thread 0!
```

Note:

- The source code for a CUDA program that prints a greeting from each thread on the GPU is shown in Program
- As you might guess, there's a header file for CUDA programs, which we include in Line 2
- The Hello function follows the include directives and starts on Line 5.
 - o This function is run by each thread on the GPU.
 - o In CUDA parlance, it's called a **kernel**, a function that is started by the host but runs on the device.
 - CUDA kernels are identified by the keyword `__global__`, and they always have return type `void`.
- The main function follows the kernel on Line 12.
 - o Like ordinary C programs, CUDA C programs start execution in main, and the main function runs on the host.
 - o The function first gets the number of threads from the command line.
 - o It then starts the required number of copies of the kernel on Line 18.
 - o The call to `cudaDeviceSynchronize` will cause the main program to wait until all the threads have finished executing the kernel, and then this

- Output

```
Hello from thread 0!
```

If we want to run ten threads on the GPU, we can type

```
$ ./cuda_hello 10
```

and the output of will be

```
Hello from thread 0!  
Hello from thread 1!  
Hello from thread 2!  
Hello from thread 3!  
Hello from thread 4!  
Hello from thread 5!  
Hello from thread 6!  
Hello from thread 7!  
Hello from thread 8!  
Hello from thread 9!
```

- The call to `cudaDeviceSynchronize` will cause the main program to wait until all the threads have finished executing the kernel, and when this happens, the program terminates as usual with return 0.

- How to compile:

- A CUDA program file that contains both host code and device code should be stored in a file with a “.cu” suffix.
 - For example, our hello program is in a file called `cuda_hello.cu`.
- We can compile it using the CUDA compiler `nvcc`.

▪ Command:

```
$ nvcc -o cuda_hello cuda_hello.cu
```

- Running it:

If we want to run one thread on the GPU, we can type

```
$ ./cuda_hello 1
```

and the output will be

So what exactly happens when we run `cuda_hello`?

- As we noted earlier, execution begins on the host in the main function. It gets the number of threads from the command line by calling the C library `strtol` function
- On line 18, the `Hello <<<1,thread_count>>>()`;
 - We tell the system how many threads to start on the GPU.
 - If there were any arguments in the `Hello` function we would put them in the parenthesis
- The kernel (which in this case is the `Hello` function denoted by `__global__`) specifies the code that each thread will execute. So each of our threads will print a message using `printf("Hello from thread %d\n, threadIdx.x");`
 - The `threadIdx.x` is a struct variable provided by CUDA when a kernel is started
 - The field `x` gives us the rank
- After a thread has printed its message, it terminates execution.
- Notice that our kernel code uses the Single-Program Multiple-Data or SPMD paradigm: each thread runs a copy of the same code on its own data. In this case, the only thread-specific data is the thread rank stored in `threadIdx.x`.
- One very important difference between the execution of an ordinary C function and a CUDA kernel is that kernel execution is asynchronous.
 - What we mean by asynchronous is that once the host (cpu) makes the call to the kernel that it should start running it will return even though the threads executing the kernel may not have finished executing.
 - Essentially all we do from the host side is make the call to start running and then we return to main while the threads run the kernels
- The call to `cudaDeviceSynchronize` in Line 21 forces the main function to wait until all the threads executing the kernel have completed.
 - If we don't use this here the program could terminate before we even get any output from the threads
- In summary:

- Execution begins in `main`, which is running on the host.
- The number of threads is taken from the command line.
- The call to `Hello` starts the kernel.
 - The `<<<1, thread_count>>>` in the call specifies that `thread_count` copies of the kernel should be started on the device.
- - When the kernel is started, the struct `threadIdx` is initialized by the system, and in our example the field `threadIdx.x` contains the thread's index or rank.
 - Each thread prints its message and terminates.
- The call to `cudaDeviceSynchronize` in `main` forces the host to wait until all of the threads have completed kernel execution before continuing and terminating.

Threads, blocks, and grids

- You're probably wondering why we put a "1" in the angle brackets in our call to `Hello`:
 - o `Hello <<<1, thread_count >>>()`;
- Recall that an Nvidia GPU consists of a collection of streaming multiprocessors (SMs), and each streaming multiprocessor consists of a collection of streaming processors (SPs). When a CUDA kernel runs, each individual thread will execute its code on an SP.
- With the "1" as the first value in the brackets we specify that all the threads started by the kernel call will run on one SM.
 - o If our GPU had two SM's we can try to use both of them with the kernel call:
`Hello <<<2, thread_count/2 >>>()`;
 - o If `thread_count` is even, this kernel call will start a total of `thread_count` threads, and the threads will be divided between the two SMs: `thread_count/2` threads will run on each SM.
- CUDA organizes threads into blocks and grids.
 - o Thread block (block) is a collection of threads that run on a single SM
 - In a kernel call, like `Hello<<<1, thread_count>>>()` the first value in the angled brackets specifies the number of thread blocks we want.
 - The second value `thread_count` is the number of threads we want in each block
 - o `Hello<<<1, thread_count>>>()` means one thread block, which consist of `thread_count` threads, and, as a consequence, we only used one SM
- We can modify our greetings program so that it uses a user-specified number of blocks, each consisting of a user-specified number of threads.

```

1  #include <stdio.h>
2  #include <cuda.h>  /* Header file for CUDA */
3
4  /* Device code: runs on GPU */
5  __global__ void Hello(void) {
6
7      printf("Hello from thread %d in block %d\n",
8             threadIdx.x, blockIdx.x);
9  } /* Hello */
10
11
12 /* Host code: Runs on CPU */
13 int main(int argc, char* argv[]) {
14     int blk_ct;          /* Number of thread blocks */
15     int th_per_blk;      /* Number of threads in each block */
16
17     blk_ct = strtol(argv[1], NULL, 10);
18     /* Get number of blocks from command line */
19     th_per_blk = strtol(argv[2], NULL, 10);
20     /* Get number of threads per block from command line */
21
22     Hello <<<blk_ct, th_per_blk>>>();
23     /* Start blk_ct*th_per_blk threads on GPU, */
24
25     cudaDeviceSynchronize(); /* Wait for GPU to finish */
26
27     return 0;
28 } /* main */

```

Program 6.2: CUDA program that prints greetings from threads in multiple blocks.

In this program we get both the number of thread blocks and the number of threads in each block from the command line. Now the kernel call starts `blk_ct` thread blocks, each of which contains `th_per_blk` threads.

When the kernel is started, each block is assigned to an SM, and the threads in the block are then run on that SM. The output is similar to the output from the original program, except that now we're using two system-defined variables: `threadIdx.x` and `blockIdx.x`.

As you've probably guessed, `threadIdx.x` gives a thread's rank or index in its block, and `blockIdx.x` gives a block's rank in the grid.

- A grid is just a collection of thread blocks started by a kernel

- A thread block consists of threads and a grid consists of thread blocks
- There are several built-in variables that a thread can use to get information on the grid started by the kernel.
 - threadIdx: the rank or index of the thread in its thread block.
 - blockDim: the dimensions, shape, or size of the thread blocks.
 - blockIdx: the rank or index of the block within the grid
 - gridDim: the dimensions, shape, or size of the grid.
- All of these structs have three fields, x, y, and z, 5 and the fields all have unsigned integer types.
 - The fields are often convenient for applications.
 - For example, an application that uses graphics may find it convenient to assign a thread to a point in two- or
 - three-dimensional space, and the fields in threadIdx can be used to indicate the point's position.

When we call a kernel with something like

```
int blk_ct, th_per_blk;
...
Hello <<<blk_ct, th_per_blk>>>();
```

the three-element structures `gridDim` and `blockDim` are initialized by assigning the values in angle brackets to the `x` fields. So, effectively, the following assignments are made:

```
gridDim.x = blk_ct;
blockDim.x = th_per_blk;
```

The `y` and `z` fields are initialized to 1. If we want to use values other than 1 for the `y` and `z` fields, we should declare two variables of type `dim3`, and pass them into the call to the kernel. For example,

```
dim3 grid_dims, block_dims;
grid_dims.x = 2;
grid_dims.y = 3;
grid_dims.z = 1;
block_dims.x = 4;
block_dims.y = 4;
block_dims.z = 4;
...
Kernel <<<grid_dims, block_dims>>> (...);
```

This should start a grid with $2 \times 3 \times 1 = 6$ blocks, each of which has $4^3 = 64$ threads.

Note that all the blocks must have the same dimensions. More importantly, CUDA requires that thread blocks be independent. So one thread block must be able to complete its execution, regardless of the states of the other thread blocks: the thread blocks can be executed sequentially in any order, or they can be executed in parallel. This ensures that the GPU can schedule a block to execute solely on the basis of the state of that block: it doesn't need to check on the state of any other block.⁶

Note:

- If we want to change the dimensions of the grid and blocks we must do so manually
 - We must declare two variables with the types `dim3` (given by CUDA) and pass them into the call to kernel (given by CUDA)
- Also note that all the blocks, in this case we are referring to the code `block_dims.[]` must have the same dimensions, aka same number value